

# Topic 12: Vector Database

1. What is a Vector Database? An easy-to-follow introduction to vector databases, explaining how they store AI embeddings for fast similarity search:  
<https://www.pinecone.io/learn/series/faiss/vector-indexes/>
2. Beginner-friendly overview of vector databases:  
<https://weaviate.io/blog/what-is-a-vector-database>
3. What is a Vector Database? A high-level video introduction that demystifies vector databases: <https://www.youtube.com/watch?v=He4Ft2lzNwc>
4. A comprehensive overview of vector search indexes:  
<https://www.pinecone.io/learn/series/faiss/vector-indexes/>
5. An insight into the production challenges of vector search. Explains why simply plugging in a “great” ANN algorithm like HNSW isn’t sufficient for scalable systems:  
<https://www.pinecone.io/learn/series/faiss/vector-indexes/>
6. A Complete Guide to Filtering in Vector Search: A detailed tutorial on using metadata filters to refine vector search results: <https://qdrant.tech/articles/vector-search-filtering/>
7. Google’s ScaNN: <https://github.com/google-research/google-research/tree/master/scann>
8. Meta’s FAISS: <https://github.com/facebookresearch/faiss>